

---

# Sparse Feature Preservation: Formally Verified Forgetting Mitigation via Low-Dimensional Activation Drift

---

Georgios Konstantopoulos  
Paradigm  
georgios@paradigm.xyz

## Abstract

Sequential fine-tuning of large language models causes capability regressions on previously learned tasks. We propose **Sparse Feature Preservation** (SFP), a regularization method that penalizes activation drift in a small, empirically identified subspace of important feature dimensions. SFP is motivated by a falsifiable hypothesis (H1): forgetting is predictable from drift in a few ablation-ranked activation directions rather than from total representational change. We formalize the mathematical core of SFP in Lean 4 with zero sorry axioms. The central result establishes a conditional reduction: if H1 holds and the preservation loss is bounded by  $\delta$ , then forgetting is bounded by  $C\sqrt{\delta}$ . Experiments on Qwen2.5-1.5B in a Math→Code setting show that top-ranked dimensions predict forgetting with  $R^2 = 0.933$ , outperforming random (0.850) and bottom-ranked (0.827) dimensions. A causal intervention test confirms this at 1.5B: perturbing important dimensions at inference causes  $2.5\times$  more degradation than random and  $28\times$  more than bottom-ranked dimensions. We report both positive and negative results: H1 does not hold at 135M parameters under our protocol but holds at 1.5B, and the margin over random controls is consistent but modest.

## 1 Introduction

Large language models are increasingly fine-tuned for specialized tasks—mathematics, code generation, safety alignment—yet each round of fine-tuning risks degrading performance on previously learned capabilities. This phenomenon, *catastrophic forgetting* [McCloskey and Cohen, 1989, French, 1999], remains a central obstacle to continual learning in LLMs [De Lange et al., 2021]. Existing mitigations include replay buffers [Chaudhry et al., 2019b], knowledge distillation [Li and Hoiem, 2017], parameter regularization [Kirkpatrick et al., 2017], and architectural constraints [Qin et al., 2024], but few provide a formally verified link between a measurable preserved quantity and a bound on forgetting.

We propose **Sparse Feature Preservation** (SFP), a method grounded in a specific mechanistic hypothesis: that old-task performance degradation is largely explained by drift in a small number of activation feature dimensions, rather than by diffuse changes across the full activation space. If this hypothesis holds, then preserving only the important feature dimensions during training should suffice to prevent forgetting, while leaving the model free to learn in the complementary subspace.

Our approach proceeds in three stages. First, we extract a PCA basis from mid-layer activations on a small memory set of old-task examples; we use PCA to obtain an orthonormal coordinate system that enables controlled comparisons across dimension subsets. Second, we rank feature dimensions by *ablation importance*: the increase in old-task loss when each dimension is zeroed out. Third, we add a preservation loss that penalizes drift in the top- $r$  most important dimensions during fine-tuning.

This selective preservation is more targeted than full hidden-state distillation [Douillard et al., 2020], which resists all representational change, and differs from replay in that it identifies *which* features to preserve rather than reusing old examples during training.

A distinctive aspect of this work is that we formally verify the mathematical properties of SFP in Lean 4 using Mathlib [The mathlib Community, 2020]. The central result, `sfp_reduces_forgetting`, establishes a conditional reduction: *assuming* our empirical hypothesis H1, bounding the preservation loss by  $\delta$  implies forgetting is bounded by  $C\sqrt{\delta}$ . All auxiliary lemmas—projection properties, loss non-negativity, gradient orthogonality at stationarity, Pythagorean decomposition of drift—are machine-checked with zero `sorry` axioms. This makes the single unverified assumption explicit: the empirical hypothesis H1 itself.

We test H1 on two model scales (SmolLM2-135M and Qwen2.5-1.5B) in a Math→Code continual learning setting. At 1.5B parameters, drift in the top- $r$  PCA dimensions achieves  $R^2 = 0.933$  for predicting forgetting, compared to 0.850 for random and 0.827 for bottom-ranked dimensions—the correct ranking predicted by H1, though with modest margins. A causal intervention test provides complementary evidence: perturbing top- $r$  dimensions at inference causes substantially more performance degradation ( $\Delta\text{ppl} = +0.85$ ) than perturbing random (+0.21) or bottom (+0.02) dimensions, even at 135M where the correlational H1 test does not hold. We report this scale dependence transparently.

### Contributions.

1. A **falsifiable measurement protocol** for the low-dimensional forgetting hypothesis, with pre-registered ablation controls (random basis, random selection, bottom- $r$ ).
2. **Sparse Feature Preservation**, a continual learning method that selectively preserves ablation-important activation features ( $\sim 40$  lines of code), with a gradient-informed variant.
3. **Formal verification** of SFP’s mathematical reduction in Lean 4: all theorems machine-checked with zero `sorry`, making assumptions explicit.
4. **Empirical evaluation** at two model scales (135M–1.5B) with honest reporting of both positive and negative results, including scale-dependent effects.

## 2 Related Work

**Continual learning and catastrophic forgetting.** Catastrophic forgetting—the tendency of neural networks to lose previously learned knowledge when trained on new tasks—has been studied since McCloskey and Cohen [1989] and French [1999]. The continual learning literature organizes mitigations into regularization-based, replay-based, and architecture-based families [De Lange et al., 2021]. SFP draws from all three: it regularizes activations (not weights), uses a small memory buffer for basis construction, and operates within a fixed architecture.

**Regularization methods.** Elastic Weight Consolidation [EWC; Kirkpatrick et al., 2017] penalizes changes to parameters deemed important via the diagonal Fisher information matrix. Synaptic Intelligence [Zenke et al., 2017] computes importance online from the path integral of gradients. These methods operate in parameter space. SFP instead operates on a low-rank subspace of the activation space, which is more directly tied to model behavior: two parameter configurations can produce identical activations, but identical activations produce identical outputs.

**Knowledge distillation.** Learning without Forgetting [LwF; Li and Hoiem, 2017] distills the old model’s output logits into the new model. PODNet [Douillard et al., 2020] extends distillation to intermediate feature maps, penalizing changes in pooled hidden-state statistics across all dimensions. SFP is related to PODNet in that both preserve intermediate representations, but SFP is *selective*: it penalizes drift in only the  $r$  dimensions that matter most for old-task performance (typically  $r = 32$  out of  $d = 1536$ ), leaving the complementary subspace free for new-task learning.

**Replay and memory.** Experience Replay [Chaudhry et al., 2019b] maintains a small buffer of old-task examples and interleaves them during training. A-GEM [Chaudhry et al., 2019a] uses the buffer to project gradients away from directions that increase old-task loss. SFP uses a memory buffer

only for basis construction—extracting PCA directions and computing ablation importance—not for direct training. The buffer is used once to build the basis, then discarded.

**Architectural methods.** Progressive Neural Networks [Rusu et al., 2016] allocate new parameters per task. Orthogonal LoRA methods [Qin et al., 2024] constrain low-rank adapter updates to be orthogonal to the subspace of previous adapters. SFP shares the geometric intuition—learning should occur in a subspace complementary to important features—but identifies the important subspace empirically via ablation rather than imposing it through adapter orthogonality.

**Mechanistic interpretability.** Recent work on sparse autoencoders [Bricken et al., 2023, Cunningham et al., 2024] and activation patching [Meng et al., 2022, Wang et al., 2023] has shown that individual activation directions can encode interpretable, causally relevant features. SFP’s ablation importance ranking is conceptually related to activation patching: both assess the causal role of activation directions by measuring downstream effects of interventions.

**Formal verification in machine learning.** Formal verification of neural network properties has focused on robustness certification [Katz et al., 2017, Wong and Kolter, 2018]. To our knowledge, no prior continual learning method has its core mathematical reduction fully machine-checked in a proof assistant. Our Lean 4 formalization [The mathlib Community, 2020] verifies the complete chain from the empirical hypothesis to a forgetting bound, making the logical structure transparent and the assumptions explicit.

### 3 Method

We describe Sparse Feature Preservation in three parts: the empirical hypothesis that motivates the method (section 3.1), the preservation algorithm (section 3.2), and the formally verified mathematical properties (section 3.3).

#### 3.1 Hypothesis: Forgetting is Low-Dimensional

Let  $f_\theta$  be a language model with parameters  $\theta$ , and let  $a_\ell(\theta) \in \mathbb{R}^d$  denote the mean-pooled activation vector at layer  $\ell$  for a given input. When we fine-tune from  $\theta^*$  (anchor) to  $\theta'$  (fine-tuned), the activation drift at layer  $\ell$  is  $\Delta a_\ell = a_\ell(\theta') - a_\ell(\theta^*)$ .

Let  $U \in \mathbb{R}^{d \times k}$  be a PCA basis computed from old-task activations at layer  $\ell$ , and let  $U_r \in \mathbb{R}^{d \times r}$  ( $r \ll k \ll d$ ) be the submatrix of the top- $r$  most important dimensions ranked by ablation. The projected drift in the important subspace is  $P_W(\Delta a_\ell) = U_r U_r^\top \Delta a_\ell$ , where  $W = \text{col}(U_r)$ .

**Hypothesis 3.1** (H1: Low-Dimensional Forgetting). *There exists a constant  $C > 0$  such that old-task performance degradation satisfies*

$$\text{forgetting} \leq C \cdot \|P_W(\Delta a)\|, \quad (1)$$

*and the projected drift  $\|P_W(\Delta a)\|$  using the top- $r$  importance-ranked dimensions is a better predictor of forgetting than total drift  $\|\Delta a\|$ , random-dimension drift, or bottom-ranked drift.*

This hypothesis is *falsifiable*: we pre-register four predictors (top- $r$ , total, random, bottom- $r$ ) and require the ranking top  $>$  total  $>$  random  $>$  bottom in  $R^2$  for predicting forgetting across training checkpoints.

**Ablation importance.** Given the PCA basis  $U = [u_1, \dots, u_k]$ , the importance of dimension  $i$  is the increase in old-task loss when that dimension is zeroed out:

$$\text{imp}(i) = \mathcal{L}_{\text{old}}(a_\ell - (u_i^\top a_\ell) u_i) - \mathcal{L}_{\text{old}}(a_\ell), \quad (2)$$

averaged over the memory set. The top- $r$  dimensions are those with the largest importance scores.

#### 3.2 SFP Algorithm

SFP operates in two phases: *basis construction* (before training) and *preservation* (during training).

---

**Algorithm 1** Sparse Feature Preservation

---

**Require:** Model  $f_{\theta^*}$ , memory set  $\mathcal{M}$ , new data  $\mathcal{D}$ , layers  $\mathcal{L}$ , dims  $k, r$ , weight  $\lambda$

```
1: // Phase 1: Basis construction
2: for  $\ell \in \mathcal{L}$  do
3:   Collect activations  $A_\ell = \{a_\ell(\theta^*, x) : x \in \mathcal{M}\}$ 
4:    $U_\ell \leftarrow \text{PCA}(A_\ell, k)$ ; rank dims by ablation importance
5:    $U_{r,\ell} \leftarrow$  top- $r$  columns; store  $a_\ell^*$ 
6: end for
7: // Phase 2: Fine-tuning
8: for each batch  $B \sim \mathcal{D}$  do
9:    $\mathcal{L} \leftarrow \mathcal{L}_{\text{new}}(B; \theta) + \lambda \sum_{\ell} \|U_{r,\ell}^\top a_\ell(\theta) - U_{r,\ell}^\top a_\ell^*\|^2$ 
10:  Update  $\theta$  via gradient descent
11: end for
```

---

**Phase 1: Basis construction.** Given a memory set  $\mathcal{M}$  of old-task examples (typically 128), we: (1) collect activations at probe layers  $\mathcal{L} = \{L/4, L/2, 3L/4\}$ ; (2) compute PCA basis  $U_\ell \in \mathbb{R}^{d \times k}$  per layer; (3) compute ablation importance per dimension; (4) select top- $r$  dimensions, yielding  $U_{r,\ell}$ ; (5) record reference activations  $a_\ell^* = a_\ell(\theta^*)$ .

**Phase 2: Preservation during training.** The training loss combines the new-task objective with a preservation penalty:

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{new}}(\theta) + \lambda \sum_{\ell \in \mathcal{L}} \underbrace{\|U_{r,\ell}^\top a_\ell(\theta) - U_{r,\ell}^\top a_\ell(\theta^*)\|^2}_{\mathcal{L}_{\text{pres}, \ell}(\theta)}, \quad (3)$$

where  $\lambda > 0$  controls the stability–plasticity trade-off. By construction, the model is free to change activations in the complement  $W_\ell^\perp$ .

**Gradient-informed variant.** As an alternative to PCA, we compute  $G = \partial \mathcal{L}_{\text{old}} / \partial a_\ell$  on the memory set and take the top- $r$  right singular vectors of  $G$  as the preservation basis. This directly identifies directions the old-task loss is most sensitive to.

### 3.3 Formally Verified Properties

We formalize SFP’s mathematical properties in Lean 4 using Mathlib (v4.17.0). The formalization comprises approximately 1,200 lines of Lean with **zero** sorry axioms.

**Definition 3.2** (Preservation Loss, `preservationLoss`). For an orthogonal projection  $P_W$  onto subspace  $W$ ,  $\mathcal{L}_{\text{pres}} = \|P_W(\Delta a)\|^2$ .

**Proposition 3.3** (`preservationLoss_nonneg`, `preservationLoss_eq_zero_iff`).  $\mathcal{L}_{\text{pres}} \geq 0$ , with equality iff  $\Delta a \in W^\perp$ .

**Proposition 3.4** (Pythagorean Decomposition, `drift_pythagorean_decomposition`). Total drift decomposes as  $\|\Delta a\|^2 = \mathcal{L}_{\text{pres}} + \|P_{W^\perp}(\Delta a)\|^2$ . Hence  $\mathcal{L}_{\text{pres}} \leq \|\Delta a\|^2$ : preserving the important subspace is strictly easier than preventing all drift.

**Proposition 3.5** (Gradient Orthogonality, `gradient_orthogonality`). At a stationary point of  $\mathcal{L}_{\text{new}} + \lambda \mathcal{L}_{\text{pres}}$  where  $\mathcal{L}_{\text{pres}} = 0$ , the new-task gradient satisfies  $P_W(\nabla \mathcal{L}_{\text{new}}) = 0$ , i.e.  $\nabla \mathcal{L}_{\text{new}} \in W^\perp$ . New-task learning occurs only in directions orthogonal to the preserved features.

**Theorem 3.6** (Main Reduction, `sfp_reduces_forgetting`). Assume H1 (theorem 3.1): forgetting  $\leq C \cdot \|P_W(\Delta a)\|$  for  $C \geq 0$ . If  $\mathcal{L}_{\text{pres}} \leq \delta$ , then

$$\text{forgetting} \leq C\sqrt{\delta}. \quad (4)$$

*Proof sketch.*  $\mathcal{L}_{\text{pres}} = \|P_W(\Delta a)\|^2 \leq \delta$  implies  $\|P_W(\Delta a)\| \leq \sqrt{\delta}$ . By H1, forgetting  $\leq C\sqrt{\delta}$ . The full Lean proof verifies that  $P_W$  is an orthogonal projection (idempotent, self-adjoint, norm-contractive).  $\square$

Table 1:  $R^2$  of linear regression predicting forgetting from activation drift (Qwen2.5-1.5B, Math→Code, 500 steps, 20 checkpoints). H1 predicts: top > total > random > bottom.

| Predictor              | Dims | $R^2$        |
|------------------------|------|--------------|
| Top- $r$ (ablation)    | 32   | <b>0.933</b> |
| Total drift            | 128  | 0.888        |
| Random dims            | 32   | 0.850        |
| Bottom- $r$ (ablation) | 32   | 0.827        |

**Theorem 3.7** (Stability–Plasticity Trade-off, `stability_plasticity_tradeoff`). *Under the same conditions, (i) forgetting  $\leq C\sqrt{\delta}$  (stability), and (ii)  $\|P_{W^\perp}(\Delta a)\|^2 \geq \|\Delta a\|^2 - \delta$  (plasticity). As  $\lambda \rightarrow \infty$ , forgetting vanishes but learning is confined to  $W^\perp$ .*

**Scope of verification.** The formalization verifies the mathematical chain from H1 to the forgetting bound. It does *not* verify H1 itself—that is an empirical claim tested in section 4. This separation is deliberate: the verification makes the logical structure transparent, so the reader knows exactly which claim requires empirical evidence.

## 4 Experiments

We evaluate two claims: (1) H1—that forgetting is predicted by drift in importance-ranked activation dimensions (section 4.2), and (2) that top-ranked dimensions are causally important for old-task performance (section 4.3). We also report a dimensionality sweep (section 4.4).

### 4.1 Setup

**Models.** We evaluate at two scales: SmolLM2-135M-Instruct (135M parameters, CPU) and Qwen2.5-1.5B-Instruct (1.5B parameters, A10G GPU). Both are fine-tuned with LoRA [Hu et al., 2022] (rank 32,  $\alpha = 64$ , targets: `q_proj`, `v_proj`).

**Tasks.** Our primary setting is Math→Code: the model is first evaluated on GSM8K [Cobbe et al., 2021], then fine-tuned on MBPP [Austin et al., 2021] for 500 steps with checkpoints every 25 steps (20 evaluation points). Forgetting is measured as the increase in math perplexity after code fine-tuning.

**SFP configuration.** Probe layers at  $L/4$ ,  $L/2$ ,  $3L/4$ . PCA with  $k = 128$  components, top  $r = 32$  dimensions by ablation importance, memory set of 128 examples, seed 42.

**Predictors.** We compare four subspaces as forgetting predictors: **top- $r$**  (ablation-ranked), **total** (all 128 dims), **random** (uniformly sampled  $r$  dims), and **bottom- $r$**  (least important).  $R^2$  is computed via linear regression of forgetting on projected drift across checkpoints.

### 4.2 H1 Validation: Correlational Evidence

**1.5B results.** Table 1 reports  $R^2$  values at the 1.5B scale. Top- $r$  achieves  $R^2 = 0.933$ , outperforming total (0.888), random (0.850), and bottom (0.827). The ranking matches H1’s prediction exactly: top > total > random > bottom.

Against our pre-registered gate: (i)  $R^2(\text{top-}r) = 0.933 > 0.6 \checkmark$ ; (ii)  $R^2(\text{top-}r) > R^2(\text{total}) \checkmark$ ; (iii)  $R^2(\text{top-}r) \gg R^2(\text{random})$ —the gap is 0.083, which is in the correct direction but modest ( $\times$ ). We score this as a **partial pass** (2/3).

**135M results.** At 135M, H1 does not hold: the ranking is inverted, with bottom-ranked dimensions ( $R^2 = 0.994$ ) outperforming top-ranked (0.913). We do not have a confirmed explanation for this inversion. Possible contributing factors include the small activation dimension ( $d = 576$ ), minimal PCA ( $k = 16$ ,  $r = 8$ ), and only 5 data points from 40 training steps. This negative result highlights that H1 is scale-dependent under our protocol.

Table 2: Causal intervention: perplexity increase ( $\Delta\text{ppl}$ ) when injecting noise into different dimension sets at inference time (Qwen2.5-1.5B, Math task,  $r = 32$ ,  $k = 128$ ). Baseline ppl = 0.505. Higher  $\Delta\text{ppl}$  = greater causal importance. Top- $r$  consistently causes the most degradation across all noise scales.

| Noise scale | Top- $r$ $\Delta\text{ppl}$ | Random $\Delta\text{ppl}$ | Bottom- $r$ $\Delta\text{ppl}$ |
|-------------|-----------------------------|---------------------------|--------------------------------|
| $0.5\sigma$ | +0.042                      | +0.024                    | +0.001                         |
| $1.0\sigma$ | +0.248                      | +0.100                    | +0.009                         |
| $2.0\sigma$ | +2.604                      | +0.810                    | +0.041                         |
| $5.0\sigma$ | +8.709                      | +5.041                    | +0.328                         |

**Baseline math performance.** The pre-fine-tuning math accuracy for Qwen2.5-1.5B is 44.44% on GSM8K (9 test questions), with baseline perplexity 0.505. SmoLLM2-135M achieves 0% accuracy (baseline perplexity 1.177).

### 4.3 H1 Validation: Causal Evidence

Correlation does not establish causation. To test whether top- $r$  dimensions are *causally* important for old-task performance, we inject calibrated Gaussian noise into specific activation subspaces at inference time and measure the resulting perplexity increase.

**Protocol.** For each dimension set (top- $r$ , random, bottom- $r$ ), we add noise  $\epsilon \sim \mathcal{N}(0, \sigma^2 I_r)$  to the corresponding PCA components at each probe layer during evaluation.  $\sigma$  is calibrated to the mean standard deviation of the projected activations, so perturbation magnitudes are comparable across dimension sets.

**Results.** Table 2 reports results at both scales across four noise magnitudes.

At  $1\sigma$ , perturbing top- $r$  dimensions causes  $2.5\times$  more degradation than random and  $28\times$  more than bottom-ranked dimensions. At  $2\sigma$  the ratio reaches  $3.2\times$  over random and  $64\times$  over bottom. The effect is monotonically increasing with noise scale and the ranking top  $>$  random  $>$  bottom holds at every scale tested.

**Cross-scale comparison.** The causal signal is present at both 135M and 1.5B. At 135M ( $r = 2$ ,  $k = 4$ ,  $1\sigma$ ), top- $r$  perturbation causes  $\Delta\text{ppl} = +0.85$  versus  $+0.21$  (random) and  $+0.02$  (bottom)—a  $4\times$  and  $42\times$  advantage respectively. This is notable because the *correlational* H1 test fails at 135M (the ranking is inverted), yet the causal test still correctly identifies the ablation-ranked dimensions as most important. The discrepancy suggests that ablation importance identifies causally important dimensions even at small scale, but the correlation between drift in those dimensions and forgetting is confounded when model capacity is limited.

### 4.4 Dimensionality Sweep

To characterize the effective dimensionality of forgetting, we sweep  $r \in \{1, 2, 4, 8, 16, 32, 64, 128\}$  at 1.5B (table 3).

Two observations: (1) a single PCA dimension captures 96.2% of activation variance, consistent with the concentrated spectrum of transformer activations; (2) forgetting  $R^2$  increases only modestly from  $r = 1$  to  $r = 128$ , suggesting the forgetting-relevant signal is concentrated in very few dimensions. The ablation-ranked  $R^2$  in table 1 (0.933 for  $r = 32$ ) exceeds the PCA-ranked value here (0.882), confirming that ablation importance provides a better ranking than variance explained alone.

### 4.5 Full Benchmark

[TODO: Five-task benchmark (Math, Code, Instruction Following, Safety, Domain) with all baselines (Naive, Replay, LwF, Hidden Distillation, Orthogonal LoRA, SFP, SFP-random-basis, SFP-random-selection) pending GPU runs. Will report Average Accuracy, Backward Transfer, Forward Transfer, and per-task forgetting at both scales.]

Table 3:  $r$ -sweep (1.5B): variance explained and forgetting  $R^2$  by top- $r$  PCA dimensions. Even  $r = 1$  captures 96.2% of variance and  $R^2 = 0.854$  of forgetting.

| $r$ | Var. Explained (%) | Forgetting $R^2$ |
|-----|--------------------|------------------|
| 1   | 96.2               | 0.854            |
| 2   | 96.4               | 0.854            |
| 4   | 96.9               | 0.871            |
| 8   | 97.4               | 0.878            |
| 16  | 98.0               | 0.879            |
| 32  | 98.7               | 0.882            |
| 64  | 99.4               | 0.883            |
| 128 | 100.0              | 0.884            |

## 5 Conclusion

We presented Sparse Feature Preservation (SFP), a method for mitigating catastrophic forgetting in LLMs by penalizing activation drift in a small, empirically identified subspace of important feature dimensions.

**Summary.** At 1.5B parameters, drift in top- $r$  ablation-ranked dimensions predicts forgetting with  $R^2 = 0.933$ , outperforming random (0.850) and bottom-ranked (0.827) dimensions. A causal intervention test confirms that ablation importance identifies causally important features: perturbing top- $r$  dimensions causes  $4\times$  more degradation than random dimensions, even at 135M where the correlational test does not hold. The Lean 4 formalization establishes a clean conditional reduction—assuming H1, bounding  $\mathcal{L}_{\text{pres}} \leq \delta$  implies forgetting  $\leq C\sqrt{\delta}$ —with zero `sorry` axioms.

### Limitations.

1. **Scale.** H1 does not hold at 135M under our protocol. We have not tested at 7B+.
2. **LoRA confound.** All experiments use LoRA, which constrains updates to a low-rank subspace. This may artificially reduce the dimensionality of drift. Full fine-tuning experiments are needed.
3. **The constant  $C$ .** The bound  $C\sqrt{\delta}$  depends on an empirically estimated  $C$ . If  $C$  is large, the bound is vacuous.
4. **Task coverage.** H1 validation uses one task pair (Math→Code). Different pairs may exhibit different forgetting geometries.
5. **Single seed.** Results are from a single random seed. Confidence intervals across seeds are not yet available.

**Future work.** (1) Scaling H1 to 7B+ models, where forgetting is a more pressing practical concern. (2) Systematic comparison of the gradient-informed basis (SVD of  $\partial\mathcal{L}_{\text{old}}/\partial a$ ) against PCA. (3) Feature interpretability (H3): determining whether top- $r$  preserved features are semantically coherent. (4) Multi-task sequences with cumulative preserved subspaces. (5) Combining SFP with replay for complementary stability–plasticity benefits.

## References

- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. In *arXiv preprint arXiv:2108.07732*, 2021.
- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.
- Arslan Chaudhry, Marc’Aurelio Ranzato, Marcus Rohrbach, and Mohamed Elhoseiny. Efficient lifelong learning with A-GEM. In *International Conference on Learning Representations*, 2019a.

- Arslan Chaudhry, Marcus Rohrbach, Mohamed Elhoseiny, Thalaiyasingam Ajber, Davide Alessio, Tomaso Poggio, and Marc Aurelio Ranato. On tiny episodic memories in continual learning. In *arXiv preprint arXiv:1902.10486*, 2019b.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. In *arXiv preprint arXiv:2110.14168*, 2021.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *International Conference on Learning Representations*, 2024.
- Matthias De Lange, Rahaf Aljundi, Marc Masana, Sarah Parisot, Xu Jia, Aleš Leonardis, Gregory Slabaugh, and Tinne Tuytelaars. A continual learning survey: Defying forgetting in classification tasks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(7):3366–3385, 2021.
- Arthur Douillard, Matthieu Cord, Charles Ollion, Thomas Robert, and Eduardo Valle. PODNet: Pooled outputs distillation for small-tasks incremental learning. In *European Conference on Computer Vision*, pages 86–102. Springer, 2020.
- Robert M French. Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3(4):128–135, 1999.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Guy Katz, Clark Barrett, David L Dill, Kyle Julian, and Mykel J Kochenderfer. Reluplex: An efficient SMT solver for verifying deep neural networks. In *International Conference on Computer Aided Verification*, pages 97–117, 2017.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Oriol Vinyals, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. In *Proceedings of the National Academy of Sciences*, volume 114, pages 3521–3526, 2017.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. In *IEEE Transactions on Pattern Analysis and Machine Intelligence*, volume 40, pages 2935–2947, 2017.
- Michael McCloskey and Neal J Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24:109–165, 1989.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems*, volume 35, 2022.
- Zhiyuan Qin, Jiyang Liu, Liang Wang, and Ran He. Exploring orthogonality in open world with LoRA. In *arXiv preprint arXiv:2410.03623*, 2024.
- Andrei A Rusu, Neil C Rabinowitz, Guillaume Desjardins, Hubert Sober, Koray Kavukcuoglu, and Raia Hadsell. Progressive neural networks. In *arXiv preprint arXiv:1606.04671*, 2016.
- The mathlib Community. The Lean mathematical library. *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*, pages 367–381, 2020.
- Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small. In *International Conference on Learning Representations*, 2023.
- Eric Wong and Zico Kolter. Provable defenses against adversarial examples via the convex outer adversarial polytope. In *International Conference on Machine Learning*, pages 5286–5295, 2018.
- Friedemann Zenke, Ben Poole, and Surya Ganguli. Continual learning through synaptic intelligence. In *International Conference on Machine Learning*, pages 3987–3995, 2017.

## A Lean 4 Theorem Index

All theorems are formalized in Lean 4 using Mathlib v4.17.0 with zero `sorry` axioms.



Table 4: Formally verified theorems referenced in this paper.

| Theorem                           | Statement                                                                                           | Module        |
|-----------------------------------|-----------------------------------------------------------------------------------------------------|---------------|
| <i>Projection properties</i>      |                                                                                                     |               |
| orthogonalProjection_idempotent   | $P_W^2 = P_W$                                                                                       | Projections   |
| orthogonalProjection_selfAdjoint  | $\langle Px, y \rangle = \langle x, Py \rangle$                                                     | Projections   |
| orthogonalProjection_norm_le      | $\ P_W(v)\  \leq \ v\ $                                                                             | Projections   |
| orthogonalProjection_complement   | $P_W + P_{W^\perp} = I$                                                                             | Projections   |
| <i>Loss properties</i>            |                                                                                                     |               |
| preservationLoss_nonneg           | $\mathcal{L}_{\text{pres}} \geq 0$                                                                  | Preservation  |
| preservationLoss_eq_zero_iff      | $\mathcal{L}_{\text{pres}} = 0 \iff \Delta a \in W^\perp$                                           | Preservation  |
| subspace_drift_bound              | $\ P_W(\Delta a)\  = \sqrt{\mathcal{L}_{\text{pres}}}$                                              | Preservation  |
| preservationLoss_le_total_drift   | $\mathcal{L}_{\text{pres}} \leq \ \Delta a\ ^2$                                                     | Preservation  |
| drift_pythagorean_decomposition   | $\ \Delta a\ ^2 = \mathcal{L}_{\text{pres}} + \ P_{W^\perp}(\Delta a)\ ^2$                          | Preservation  |
| <i>Gradient properties</i>        |                                                                                                     |               |
| gradient_orthogonality            | $P_W(\nabla \mathcal{L}_{\text{new}}) = 0$ at stationarity                                          | GradientOrth. |
| gradient_in_complement            | $\nabla \mathcal{L}_{\text{new}} \in W^\perp$ at stationarity                                       | GradientOrth. |
| <i>Main results</i>               |                                                                                                     |               |
| sfp_reduces_forgetting            | $H1 \wedge \mathcal{L}_{\text{pres}} \leq \delta \Rightarrow \text{forgetting} \leq C\sqrt{\delta}$ | Reduction     |
| sfp_reduces_forgetting_multilayer | Per-layer version                                                                                   | Reduction     |
| stability_plasticity_tradeoff     | Stability + plasticity bounds                                                                       | Reduction     |

### Building the formalization.

```
docker build -t sfp-lean lean-sfp/
docker run --rm sfp-lean build
grep -rn sorry lean-sfp/SFP/ # should return empty
```

## B Experimental Details

Table 5: Training hyperparameters.

| Parameter              | Value              |
|------------------------|--------------------|
| Optimizer              | AdamW              |
| Learning rate          | $2 \times 10^{-4}$ |
| Weight decay           | 0.01               |
| Batch size             | 4                  |
| Training steps         | 500                |
| Checkpoint interval    | 25 steps           |
| LoRA rank / $\alpha$   | 32 / 64            |
| LoRA targets           | q_proj, v_proj     |
| PCA components ( $k$ ) | 128                |
| Preserved dims ( $r$ ) | 32                 |
| Memory set             | 128 examples       |
| Probe layers           | $L/4, L/2, 3L/4$   |
| Seed                   | 42                 |

**SmolLM2-135M-Instruct.** 135M parameters,  $L = 30$  layers, probe layers at 7/15/22, activation dim  $d = 576$ . Baseline: 0% math accuracy, 1.177 math perplexity.

**Qwen2.5-1.5B-Instruct.** 1.5B parameters,  $L = 28$  layers, probe layers at 7/14/21, activation dim  $d = 1536$ . Loads in bfloat16; all hooks preserve original dtype. Baseline: 44.44% math accuracy, 0.505 math perplexity.

**Causal intervention details.** Noise  $\epsilon \sim \mathcal{N}(0, \sigma^2 I_r)$  injected at each probe layer.  $\sigma = s \cdot \bar{\sigma}_{\text{dim}}$  where  $s$  is the noise scale multiplier (default 1.0) and  $\bar{\sigma}_{\text{dim}}$  is the mean eigenvalue scale ( $\sqrt{\text{variance}}$ ) of the selected PCA dimensions. Noise generated on CPU, cast to model dtype before injection.

**Ablation importance cost.**  $k$  forward passes per layer  $\times$  3 layers  $\times$  20 eval examples =  $\sim 7,680$  forward passes for the 1.5B model ( $\sim 10$  min on A10G).

**Computational cost.** Full H1 validation (1.5B, 500 steps, 20 checkpoints):  $\sim 90$  minutes on a single A10G GPU. Causal intervention test:  $\sim 15$  minutes additional. Total cost per experiment: approximately \$2–3 on Modal.